

### 3\_Short\_list\_promising\_models

April 11, 2021

#### Completed by:

Name : Adarsh Ghimire

Student ID : 100058927

MSc in Electrical and Computer Engineering

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[ ]: try:
    import os
    os.chdir('drive/MyDrive/COSC 606 Machine learning/Project/notebooks/')
except:
    print("Already in working directory")
```

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
```

```
[ ]: # Training and test data were made in second notebook, and were saved to make_
    ↪future use easier
training_dataset_path = "../final_dataset/train.csv"
test_dataset_path = "../final_dataset/test.csv"
train_df = pd.read_csv(training_dataset_path)
test_df = pd.read_csv(test_dataset_path)
train_df.head()
```

```
[ ]: Report Date Route Time ... Direction Min Delay Min Gap
0 2014-01-02 505 06:31:00 ... E/B 4.0 8.0
1 2014-01-02 504 12:43:00 ... E/B 20.0 22.0
2 2014-01-02 501 14:01:00 ... W/B 13.0 19.0
3 2014-01-02 504 14:22:00 ... W/B 7.0 11.0
4 2014-01-02 504 16:42:00 ... E/B 3.0 6.0
```

[5 rows x 8 columns]

## 0.1 Data Manipulations functions

```
[ ]: # Since our target is Min Delay so the function below cleans the Min Delay rows
      ↳ if they are nan values
def clean_delay(df):
    df = df[df['Min Delay'].notna()]
    return df
```

```
[ ]: # This function checks for valid routes only
def check_route(x):
    # load the valid list of TTC Streetcar routes
    valid_routes = [501, 502, 503, 504, 505, 506, 509, 510, 511, 512, 301, 304,
                    ↳ 306, 310]

    if x in valid_routes:
        return x
    else:
        return "bad route"

# This function cleans the data based on valid routes
def clean_route(df):
    # This function takes dataframe as input
    # cleans the route column based on the validity of the route of street car
    # returns the cleaned dataframe
    df['Route'] = df['Route'].apply(lambda x: check_route(x))
    df = df[df.Route != "bad route"]
    df['Route'] = df['Route'].astype('int64')
    return df
```

```
[ ]: # This function drops the Location column
def drop_location(df):
    df = df.drop(["Location"], axis=1)
    return df
```

```
[ ]: def create_date_time_column(df):
    # This function takes dataframe, then merges the date and time
    # Then convert that column into datetime datatype
```

```

# Such that it can be further used in time series easily
try:
    new = pd.to_datetime(df["Report Date"] + " " + df["Time"], utc=True)
    df["Date Time"] = new
    df = df.drop(["Report Date", "Time"], axis=1)
    return df
except:
    return df

```

```

[ ]: # This function divides a day into different period
def day_divider(hour):
    if hour > 5 and hour < 12:
        return "morning"
    elif hour >= 12 and hour < 17:
        return "afternoon"
    elif hour >= 17 and hour < 21:
        return "evening"
    else:
        return "night"

```

```

[ ]: def clean_gap(df):
    # This function will help to clean the Min Gap column feature with training_
    ↪ data Min Gap mean value
    df["Min Gap"] = df["Min Gap"].fillna(train_df['Min Gap'].mean())
    return df

```

```

[ ]: # These function help to filter the Direction values and clean them
valid_directions = ['eb', 'wb', 'nb', 'sb', 'bw']
def check_direction(x):
    if x in valid_directions:
        return(x)
    else:
        return("bad direction")

def direction_cleanup(df):
    df['Direction'] = df['Direction'].str.lower()
    df['Direction'] = df['Direction'].str.replace('/', '')
    df['Direction'] = df['Direction'].replace({'eastbound': 'eb', 'westbound':
    ↪ 'wb', 'southbound': 'sb', 'northbound': 'nb'})
    df['Direction'] = df['Direction'].apply(lambda x: check_direction(x))
    return(df)

```

```

[ ]: def complete_cleaner(df):
    df = clean_delay(df) # drops the nan Min delay rows
    df = clean_route(df) # cleans the unwanted route from the dataset
    df = drop_location(df) # drops the location column from the dataset
    df = create_date_time_column(df) # creates Date Time column in the dataset

```

```

df["Part of Day"] = df.apply(lambda x: day_divider(x["Date Time"].
→hour),axis=1) # Creates Part of Day column in the dataset
df = clean_gap(df) # cleans gap based on the mean of gap values
df = direction_cleanup(df) # cleans the direction column to 5_
→directions(eb,wb,nb,sb,bw)
df = df[df["Direction"] != "bad direction"]
df.reset_index(inplace=True, drop=True)
df.drop(['Date Time'], axis=1, inplace=True)
return df

```

```

[ ]: # Creating a copy of training data to work on
df = train_df.copy(deep=True)

```

```

[ ]: df = complete_cleaner(df)

```

/usr/local/lib/python3.7/dist-packages/ipykernel\_launcher.py:16:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
app.launch_new_instance()
```

/usr/local/lib/python3.7/dist-packages/ipykernel\_launcher.py:18:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

[ ]: df.sample(10)

```

```

[ ]:

```

	Route	Day	Direction	Min Delay	Min Gap	Part of Day
11615	512	Monday	eb	4.0	8.0	morning
25712	509	Saturday	wb	18.0	36.0	night
61661	506	Tuesday	wb	6.0	12.0	afternoon
57689	506	Thursday	wb	5.0	10.0	evening
17206	501	Monday	eb	10.0	20.0	afternoon
57370	504	Wednesday	eb	11.0	21.0	night
27191	504	Monday	eb	4.0	8.0	morning
51570	501	Friday	eb	60.0	67.0	night
64756	509	Saturday	eb	13.0	26.0	evening
22315	504	Wednesday	wb	7.0	11.0	evening

## 1 Now doing one hot encoding for the categorical data

1. Route is categorical
2. Day is categorical
3. Direction is categorical
4. Part of Day is categorical

```
[ ]: # One hot encoder object made using training data, so that it can be used for
      ↪testing later on.
from sklearn.preprocessing import OneHotEncoder
import numpy as np
day_enc = OneHotEncoder()
route_enc = OneHotEncoder()
dir_enc = OneHotEncoder()
part_enc = OneHotEncoder()

day_enc.fit(np.array(df["Day"]).reshape(-1,1))
route_enc.fit(np.array(df["Route"]).reshape(-1,1))
dir_enc.fit(np.array(df["Direction"]).reshape(-1,1))
part_enc.fit(np.array(df["Part of Day"]).reshape(-1,1))
```

```
[ ]: OneHotEncoder(categories='auto', drop=None, dtype=<class 'numpy.float64'>,
                    handle_unknown='error', sparse=True)
```

```
[ ]: def one_hot_encoder(df):

    df[day_enc.categories_[0]] = day_enc.transform(np.array(df["Day"]).
    ↪reshape(-1,1)).toarray()
    df[route_enc.categories_[0]] = route_enc.transform(np.array(df["Route"]).
    ↪reshape(-1,1)).toarray()
    df[dir_enc.categories_[0]] = dir_enc.transform(np.array(df["Direction"]).
    ↪reshape(-1,1)).toarray()
    df[part_enc.categories_[0]] = part_enc.transform(np.array(df["Part of Day"]).
    ↪reshape(-1,1)).toarray()

    df.drop(['Day', 'Route', 'Direction', 'Part of Day'], axis=1, inplace=True)
    return df
```

```
[ ]: df = one_hot_encoder(df)
```

```
[ ]: df
```

```
[ ]:      Min Delay  Min Gap  Friday  Monday  ...  afternoon  evening  morning
night
0           4.0      8.0      0.0      0.0  ...         0.0        0.0        1.0
0.0
1          20.0     22.0      0.0      0.0  ...         1.0        0.0        0.0
```

```

0.0
2          13.0      19.0      0.0      0.0 ...          1.0      0.0      0.0
0.0
3          7.0       11.0      0.0      0.0 ...          1.0      0.0      0.0
0.0
4          3.0       6.0       0.0      0.0 ...          1.0      0.0      0.0
0.0
...          ...          ...          ...          ...          ...          ...          ...
...
71507       8.0      16.0      0.0      0.0 ...          0.0      0.0      1.0
0.0
71508       7.0      14.0      0.0      0.0 ...          0.0      0.0      1.0
0.0
71509       5.0      10.0      0.0      0.0 ...          0.0      0.0      1.0
0.0
71510       5.0      10.0      0.0      0.0 ...          0.0      0.0      1.0
0.0
71511       9.0      16.0      0.0      0.0 ...          0.0      0.0      1.0
0.0

```

[71512 rows x 32 columns]

## 1.1 Need to do feature scaling for Min Gap column feature

```

[ ]: # Feature Standardization, which means changing between -1 and 1, such that
      ↪ algorithm can converge swiftly
min_gap_train_mean = np.mean(df["Min Gap"])
min_gap_train_std = np.std(df["Min Gap"])
print(min_gap_train_mean)
print(min_gap_train_std)

def min_gap_scaler(df):
    df["Min Gap"] = (df["Min Gap"] - min_gap_train_mean) / min_gap_train_std
    return df

```

18.13560073354417

33.68768666213863

```

[ ]: df = min_gap_scaler(df)

```

```

[ ]: df.head()

```

```

[ ]:
   Min Delay  Min Gap  Friday  Monday  ...  afternoon  evening  morning  night
0         4.0 -0.300870      0.0     0.0  ...         0.0      0.0      1.0     0.0
1        20.0  0.114713      0.0     0.0  ...         1.0      0.0      0.0     0.0
2        13.0  0.025659      0.0     0.0  ...         1.0      0.0      0.0     0.0
3         7.0 -0.211816      0.0     0.0  ...         1.0      0.0      0.0     0.0

```

```
4          3.0 -0.360238      0.0      0.0 ...      1.0      0.0      0.0      0.0
```

```
[5 rows x 32 columns]
```

## 2 The data is ready for training

```
[ ]: def complete_data_preprocessing(df):  
    # This function is for test set  
    df = complete_cleaner(df)  
    df = one_hot_encoder(df) # one hot encoder based on training samples  
    df = min_gap_scaler(df) # feature scaler based on training sample  
    print("-----")  
    if df.isnull().values.any():  
        print("There are some null values")  
    else:  
        print("Data is ready for testing")  
    return df
```

```
[ ]: processed_test_data = complete_data_preprocessing(test_df)
```

```
-----  
Data is ready for testing
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:16:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
app.launch_new_instance()
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:18:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

**2.0.1 The regression results function computes the different metrics like r2 score, MAE, and RMSE.**

```
[ ]: import sklearn.metrics as metrics  
def regression_results(y_true, y_pred):  
    # Regression metrics  
    mean_absolute_error=metrics.mean_absolute_error(y_true, y_pred)  
    mse=metrics.mean_squared_error(y_true, y_pred)
```

```

r2 = metrics.r2_score(y_true, y_pred)
print("R2:", round(r2, 4))
print('MAE: ', round(mean_absolute_error,4))
print('RMSE: ', round(np.sqrt(mse),4))

```

```

[ ]: # Selected features for model training
selected_features = list(df.columns)
selected_features.remove("Min Delay")
target = "Min Delay"

```

Training and validation dataset will be used throughout the model selection and fine tuning process

*Final test set has been kept separate, it will be used only after the model is finalized.*

```

[ ]: X = df[selected_features]
y = df[target]
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.1,
↪random_state=42)

```

```

[ ]: def evaluate_model(model):
    print("Training Regression results:")
    regression_results(y_train, model.predict(X_train))
    print("\n")
    print("Validation Regression results")
    regression_results(y_val, model.predict(X_val))

```

---

Model 1 - Simple linear regression

```

[ ]: lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)
print("Linear regression results")
evaluate_model(lin_reg)

```

Linear regression results

Training Regression results:

R2: 0.6239

MAE: 3.9846

RMSE: 18.5552

Validation Regression results

R2: 0.8012

MAE: 3.799

RMSE: 11.0773

---

Model 2 - Polynomial regression of degree 2



```
[ ]: poly_2 = PolynomialFeatures(degree = 2)
X_poly_2 = poly_2.fit_transform(X)
poly_2.fit(X_poly_2, y)

X_poly_2_train, X_poly_2_val, y_train_poly_2, y_val_poly_2 = train_test_split(
    X_poly_2,
    y,
    test_size=0.1,
    random_state=42)

reg_2 = LinearRegression()
reg_2.fit(X_poly_2_train, y_train_poly_2)

print("Training polynomial Regression of degree 2 results:")
regression_results(y_train_poly_2, reg_2.predict(X_poly_2_train))
print("\n")
print("Validation polynomial Regression of degree 2 results")
regression_results(y_val_poly_2, reg_2.predict(X_poly_2_val))
```

Training polynomial Regression of degree 2 results:  
R2: 0.8298  
MAE: 2.9791  
RMSE: 12.4826

Validation polynomial Regression of degree 2 results  
R2: 0.852  
MAE: 2.9469  
RMSE: 9.558

---

Model 3 - Bayesian regression

```
[ ]: from sklearn import linear_model
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.1,
    ↪random_state=42)

BayReg = linear_model.BayesianRidge()
BayReg.fit(X_train, y_train)
print("Bayesian regression results")
evaluate_model(BayReg)
```

Bayesian regression results  
Training Regression results:  
R2: 0.6239  
MAE: 3.9854  
RMSE: 18.5552

Validation Regression results

R2: 0.8011

MAE: 3.7998

RMSE: 11.0804

---

Model 4 - SVM Regressor

```
[ ]: # It takes forever to train, wait for more than 20 minutes, but still
# no results were given so, terminated

# from sklearn.svm import SVR
# X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.1,
↳random_state=42)
# svr = SVR(kernel='poly', C=1.0, epsilon=0.2)
# svr.fit(X_train, y_train)
# print("SVM regression results")
# evaluate_model(svr)

# SVM regressor took forever to train, so terminated it in the middle of
↳training.
# SVM regressor not suitable for this task
```

---

Model 5 - Random Forest

```
[ ]: from sklearn.ensemble import RandomForestRegressor
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.1,
↳random_state=42)
rf = RandomForestRegressor(random_state=42)
rf.fit(X_train, y_train)
print("Random Forest regression results")
evaluate_model(rf)
```

Random Forest regression results

Training Regression results:

R2: 0.9382

MAE: 1.3267

RMSE: 7.523

Validation Regression results

R2: 0.7878

MAE: 2.3279

RMSE: 11.4465

---

## Model 6 - Neural Network

```
[ ]: from sklearn.neural_network import MLPRegressor
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.1,
↪random_state=42)
nn = MLPRegressor(random_state=1, max_iter=500).fit(X_train, y_train)
print("Neural network regression results")
evaluate_model(nn)
```

Neural network regression results

Training Regression results:

R2: 0.8424

MAE: 2.7425

RMSE: 12.0125

Validation Regression results

R2: 0.8636

MAE: 2.6741

RMSE: 9.177

---

## Model 7 - XGBoost

```
[ ]: from sklearn.ensemble import GradientBoostingRegressor
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.1,
↪random_state=42)
boost = GradientBoostingRegressor().fit(X_train, y_train)
print("Gradient Boosting network regression results")
evaluate_model(boost)
```

Gradient Boosting network regression results

Training Regression results:

R2: 0.876

MAE: 2.1704

RMSE: 10.6554

Validation Regression results

R2: 0.8709

MAE: 2.1503

RMSE: 8.9268

### 3 K-Fold Cross validation

#### 3.1 Cross validation function for model training

```
[ ]: def cross_validation(model, X, y, n=10): # default 10 splits
    from sklearn.model_selection import KFold
    kf = KFold(n_splits=n)
    kf.get_n_splits(X)

    train_scores = {}
    val_scores = {}
    train_scores["R2"], train_scores["MAE"], train_scores["RMSE"] = [], [], []
    val_scores["R2"], val_scores["MAE"], val_scores["RMSE"] = [], [], []
    model_count = 1
    for train_index, val_index in kf.split(X):

        X_train, X_val = X.loc[train_index], X.loc[val_index]
        y_train, y_val = y[train_index], y[val_index]
        model.fit(X_train, y_train)
        print("*****")
        print("Model {} trained".format(model_count))
        print("Training Scores")
        print("R2 Score : {} \t MAE : {} \t RMSE : {}".format(round(metrics.
↪r2_score(y_train, model.predict(X_train)),4),
                                                                round(metrics.
↪mean_absolute_error(y_train, model.predict(X_train)),4),
                                                                round(np.sqrt(metrics.
↪mean_squared_error(y_train, model.predict(X_train))),4)
                                                                ))
        print("Validation Scores")
        print("R2 Score : {} \t MAE : {} \t RMSE : {}".format(round(metrics.
↪r2_score(y_val, model.predict(X_val)),4),
                                                                round(metrics.
↪mean_absolute_error(y_val, model.predict(X_val)),4),
                                                                round(np.sqrt(metrics.
↪mean_squared_error(y_val, model.predict(X_val))),4)
                                                                ))

        train_scores["R2"].append(round(metrics.r2_score(y_train, model.
↪predict(X_train)),4))
        train_scores["MAE"].append(round(metrics.mean_absolute_error(y_train, model.
↪predict(X_train)),4))
        train_scores["RMSE"].append(round(np.sqrt(metrics.
↪mean_squared_error(y_train, model.predict(X_train))),4))

        val_scores["R2"].append(round(metrics.r2_score(y_val, model.
↪predict(X_val)),4))
```

```

        val_scores["MAE"].append(round(metrics.mean_absolute_error(y_val, model.
→predict(X_val)),4))
        val_scores["RMSE"].append(round(np.sqrt(metrics.mean_squared_error(y_val,
→model.predict(X_val))),4))

        model_count += 1

    return train_scores, val_scores

```

```

[ ]: def print_results(train_scores, val_scores):
    """
        train_scores, and val_scores should be a dictionary containing R2, MAE and
→RMSE
    """
    print("\n")
    print("Mean and standard deviation of {}-fold cross validation results".
→format(len(train_scores["R2"])))
    print("-----")
    print("Train Mean R2 Score: \t {} \t\t Validation Mean R2 Score: \t {}".
→format(round(np.mean(train_scores["R2"]),4), round(np.
→mean(val_scores["R2"]),4)))
    print("Train Std R2 Score: \t {} \t\t Validation Std R2 Score: \t {}".
→format(round(np.std(train_scores["R2"]),4), round(np.
→std(val_scores["R2"]),4)))
    print("\n")
    print("Train Mean MAE Score: \t {} \t\t Validation Mean MAE Score: \t {}".
→format(round(np.mean(train_scores["MAE"]),4), round(np.
→mean(val_scores["MAE"]),4)))
    print("Train Std MAE Score: \t {} \t\t Validation Std MAE Score: \t {}".
→format(round(np.std(train_scores["MAE"]),4), round(np.
→std(val_scores["MAE"]),4)))
    print("\n")
    print("Train Mean RMSE Score: \t {} \t\t Validation Mean RMSE Score: \t {}".
→format(round(np.mean(train_scores["RMSE"]),4), round(np.
→mean(val_scores["RMSE"]),4)))
    print("Train Std RMSE Score: \t {} \t\t Validation Std RMSE Score: \t {}".
→format(round(np.std(train_scores["RMSE"]),4), round(np.
→std(val_scores["RMSE"]),4)))

```

### 3.2 Linear regression

```

[ ]: train_scores, val_scores = cross_validation(LinearRegression(), X, y, 10)

```

\*\*\*\*\*

```

Model 1 trained
Training Scores
R2 Score : 0.6216      MAE : 3.9786      RMSE : 18.4698
Validation Scores
R2 Score : 0.7948      MAE : 3.8018      RMSE : 12.3318
*****
Model 2 trained
Training Scores
R2 Score : 0.6201      MAE : 3.9552      RMSE : 18.1851
Validation Scores
R2 Score : 0.756       MAE : 4.2323      RMSE : 15.7741
*****
Model 3 trained
Training Scores
R2 Score : 0.6266      MAE : 3.8001      RMSE : 17.9763
Validation Scores
R2 Score : 0.7073      MAE : 4.2025      RMSE : 17.6362
*****
Model 4 trained
Training Scores
R2 Score : 0.6221      MAE : 3.9856      RMSE : 18.3186
Validation Scores
R2 Score : 0.7646      MAE : 4.0355      RMSE : 14.2505
*****
Model 5 trained
Training Scores
R2 Score : 0.8285      MAE : 2.6946      RMSE : 12.3923
Validation Scores
R2 Score : -1.8484     MAE : 3.417       RMSE : 47.584
*****
Model 6 trained
Training Scores
R2 Score : 0.6222      MAE : 4.0218      RMSE : 17.8355
Validation Scores
R2 Score : 0.7146      MAE : 4.407       RMSE : 19.0997
*****
Model 7 trained
Training Scores
R2 Score : 0.6267      MAE : 4.03        RMSE : 18.7283
Validation Scores
R2 Score : 0.8398      MAE : 3.563       RMSE : 7.9691
*****
Model 8 trained
Training Scores
R2 Score : 0.6207      MAE : 4.0756      RMSE : 18.3995
Validation Scores
R2 Score : 0.7859      MAE : 3.4817      RMSE : 13.2421
*****

```

```

Model 9 trained
Training Scores
R2 Score : 0.6175      MAE : 4.0498      RMSE : 18.5795
Validation Scores
R2 Score : 0.8417      MAE : 3.6346      RMSE : 10.7612
*****
Model 10 trained
Training Scores
R2 Score : 0.5989      MAE : 4.0737      RMSE : 18.5448
Validation Scores
R2 Score : 0.8809      MAE : 4.2483      RMSE : 11.6323

```

```
[ ]: print("Linear Regression results : ")
      print_results(train_scores, val_scores)
```

Linear Regression results :

Mean and standard deviation of 10-fold cross validation results

```

-----
Train Mean R2 Score:      0.6405      Validation Mean R2 Score:
0.5237
Train Std R2 Score:      0.0631      Validation Std R2 Score:
0.7925

Train Mean MAE Score:      3.8665      Validation Mean MAE Score:
3.9024
Train Std MAE Score:      0.3979      Validation Std MAE Score:
0.3465

Train Mean RMSE Score:      17.743      Validation Mean RMSE Score:
17.0281
Train Std RMSE Score:      1.8026      Validation Std RMSE Score:
10.6511

```

### 3.3 Bayesian regression

```
[ ]: train_scores, val_scores = cross_validation(linear_model.BayesianRidge(), X, y,
      ↪10)
```

```

*****
Model 1 trained
Training Scores
R2 Score : 0.6216      MAE : 3.9792      RMSE : 18.4699
Validation Scores
R2 Score : 0.7947      MAE : 3.8023      RMSE : 12.3347

```

```

*****
Model 2 trained
Training Scores
R2 Score : 0.6201      MAE : 3.9595      RMSE : 18.1852
Validation Scores
R2 Score : 0.756       MAE : 4.2339      RMSE : 15.7742
*****
Model 3 trained
Training Scores
R2 Score : 0.6266      MAE : 3.801       RMSE : 17.9763
Validation Scores
R2 Score : 0.7073      MAE : 4.2038      RMSE : 17.6362
*****
Model 4 trained
Training Scores
R2 Score : 0.6221      MAE : 3.9865      RMSE : 18.3186
Validation Scores
R2 Score : 0.7645      MAE : 4.0355      RMSE : 14.2527
*****
Model 5 trained
Training Scores
R2 Score : 0.8285      MAE : 2.6946      RMSE : 12.3923
Validation Scores
R2 Score : -1.8479     MAE : 3.4173      RMSE : 47.5796
*****
Model 6 trained
Training Scores
R2 Score : 0.6222      MAE : 4.0227      RMSE : 17.8356
Validation Scores
R2 Score : 0.7145      MAE : 4.4076      RMSE : 19.1025
*****
Model 7 trained
Training Scores
R2 Score : 0.6267      MAE : 4.0309      RMSE : 18.7284
Validation Scores
R2 Score : 0.8397      MAE : 3.5633      RMSE : 7.9706
*****
Model 8 trained
Training Scores
R2 Score : 0.6207      MAE : 4.0764      RMSE : 18.3996
Validation Scores
R2 Score : 0.7859      MAE : 3.4828      RMSE : 13.2443
*****
Model 9 trained
Training Scores
R2 Score : 0.6175      MAE : 4.0507      RMSE : 18.5795
Validation Scores
R2 Score : 0.8417      MAE : 3.6356      RMSE : 10.7636

```



\*\*\*\*\*

Model 10 trained

Training Scores

R2 Score : 0.5989            MAE : 4.0748        RMSE : 18.5448

Validation Scores

R2 Score : 0.8808            MAE : 4.2496        RMSE : 11.635

```
[ ]: print("Bayesian Regression results : ")
      print_results(train_scores, val_scores)
```

Bayesian Regression results :

Mean and standard deviation of 10-fold cross validation results

-----  
Train Mean R2 Score:        0.6405                            Validation Mean R2 Score:  
0.5237

Train Std R2 Score:        0.0631                            Validation Std R2 Score:  
0.7923

Train Mean MAE Score:        3.8676                            Validation Mean MAE Score:  
3.9032

Train Std MAE Score:        0.3983                            Validation Std MAE Score:  
0.3467

Train Mean RMSE Score:        17.743                            Validation Mean RMSE Score:  
17.0293

Train Std RMSE Score:        1.8026                            Validation Std RMSE Score:  
10.6492

### 3.4 Random Forest regression

#### 3.4.1 With default parameters

```
[ ]: from sklearn.ensemble import RandomForestRegressor
      train_scores, val_scores = \
      ↪ cross_validation(RandomForestRegressor(random_state=42), X, y, 10)
```

\*\*\*\*\*

Model 1 trained

Training Scores

R2 Score : 0.936            MAE : 1.3269        RMSE : 7.5963

Validation Scores

R2 Score : 0.8438            MAE : 2.5018        RMSE : 10.758

\*\*\*\*\*

Model 2 trained

Training Scores  
 R2 Score : 0.9392                      MAE : 1.3003              RMSE : 7.2726  
 Validation Scores  
 R2 Score : 0.7908                      MAE : 2.7948              RMSE : 14.6076  
 \*\*\*\*\*  
 Model 3 trained  
 Training Scores  
 R2 Score : 0.9313                      MAE : 1.3428              RMSE : 7.7094  
 Validation Scores  
 R2 Score : 0.8554                      MAE : 2.4851              RMSE : 12.3964  
 \*\*\*\*\*  
 Model 4 trained  
 Training Scores  
 R2 Score : 0.942                      MAE : 1.3253              RMSE : 7.1799  
 Validation Scores  
 R2 Score : 0.8146                      MAE : 2.1795              RMSE : 12.6465  
 \*\*\*\*\*  
 Model 5 trained  
 Training Scores  
 R2 Score : 0.9398                      MAE : 1.301                RMSE : 7.3453  
 Validation Scores  
 R2 Score : 0.5226                      MAE : 2.8284              RMSE : 19.481  
 \*\*\*\*\*  
 Model 6 trained  
 Training Scores  
 R2 Score : 0.9434                      MAE : 1.2755              RMSE : 6.9034  
 Validation Scores  
 R2 Score : 0.736                      MAE : 2.9011              RMSE : 18.3674  
 \*\*\*\*\*  
 Model 7 trained  
 Training Scores  
 R2 Score : 0.934                      MAE : 1.3915              RMSE : 7.8732  
 Validation Scores  
 R2 Score : 0.8035                      MAE : 1.9532              RMSE : 8.8249  
 \*\*\*\*\*  
 Model 8 trained  
 Training Scores  
 R2 Score : 0.9353                      MAE : 1.3866              RMSE : 7.5979  
 Validation Scores  
 R2 Score : 0.8439                      MAE : 1.6439              RMSE : 11.306  
 \*\*\*\*\*  
 Model 9 trained  
 Training Scores  
 R2 Score : 0.9309                      MAE : 1.3724              RMSE : 7.8973  
 Validation Scores  
 R2 Score : 0.8997                      MAE : 2.0552              RMSE : 8.5652  
 \*\*\*\*\*  
 Model 10 trained

```

Training Scores
R2 Score : 0.9261      MAE : 1.3753      RMSE : 7.9614
Validation Scores
R2 Score : 0.9374      MAE : 2.2813      RMSE : 8.4341

```

```
[ ]: print("Random Forest regression results")
      print_results(train_scores, val_scores)
```

Random Forest regression results

Mean and standard deviation of 10-fold cross validation results

```

-----
Train Mean R2 Score:      0.9358      Validation Mean R2 Score:
0.8048
Train Std R2 Score:      0.0052      Validation Std R2 Score:
0.1081

Train Mean MAE Score:     1.3398      Validation Mean MAE Score:
2.3624
Train Std MAE Score:      0.0384      Validation Std MAE Score:
0.3933

Train Mean RMSE Score:    7.5337      Validation Mean RMSE Score:
12.5387
Train Std RMSE Score:     0.331      Validation Std RMSE Score:      3.7095

```

## 3.5 Gradient Boosting regression

### 3.5.1 With default parameters

```
[ ]: from sklearn.ensemble import GradientBoostingRegressor
      train_scores, val_scores = cross_validation(GradientBoostingRegressor(), X, y,
      ↪10)
```

```

*****
Model 1 trained
Training Scores
R2 Score : 0.8717      MAE : 2.1268      RMSE : 10.755
Validation Scores
R2 Score : 0.8678      MAE : 2.4526      RMSE : 9.8983
*****
Model 2 trained
Training Scores
R2 Score : 0.8795      MAE : 2.0974      RMSE : 10.2405
Validation Scores

```

```

R2 Score : 0.827          MAE : 2.6866      RMSE : 13.2818
*****
Model 3 trained
Training Scores
R2 Score : 0.8629          MAE : 2.159      RMSE : 10.8944
Validation Scores
R2 Score : 0.8872          MAE : 2.5061      RMSE : 10.9496
*****
Model 4 trained
Training Scores
R2 Score : 0.8764          MAE : 2.1553      RMSE : 10.4759
Validation Scores
R2 Score : 0.8327          MAE : 2.2145      RMSE : 12.0133
*****
Model 5 trained
Training Scores
R2 Score : 0.8876          MAE : 2.1058      RMSE : 10.0339
Validation Scores
R2 Score : 0.6278          MAE : 2.6106      RMSE : 17.2006
*****
Model 6 trained
Training Scores
R2 Score : 0.8899          MAE : 2.0705      RMSE : 9.6293
Validation Scores
R2 Score : 0.7867          MAE : 2.6321      RMSE : 16.5122
*****
Model 7 trained
Training Scores
R2 Score : 0.8716          MAE : 2.2348      RMSE : 10.9837
Validation Scores
R2 Score : 0.8654          MAE : 1.8158      RMSE : 7.3045
*****
Model 8 trained
Training Scores
R2 Score : 0.8734          MAE : 2.2182      RMSE : 10.629
Validation Scores
R2 Score : 0.8581          MAE : 1.6427      RMSE : 10.7821
*****
Model 9 trained
Training Scores
R2 Score : 0.8628          MAE : 2.2012      RMSE : 11.1295
Validation Scores
R2 Score : 0.9203          MAE : 1.9805      RMSE : 7.6348
*****
Model 10 trained
Training Scores
R2 Score : 0.8565          MAE : 2.1976      RMSE : 11.0939
Validation Scores

```

R2 Score : 0.9683            MAE : 2.1003        RMSE : 6.0022

```
[ ]: print("XGBoost regression results")
     print_results(train_scores, val_scores)
```

XGBoost regression results

Mean and standard deviation of 10-fold cross validation results

```
-----
Train Mean R2 Score:      0.8732                Validation Mean R2 Score:
0.8441
Train Std R2 Score:       0.0102                Validation Std R2 Score:
0.0865
```

```
Train Mean MAE Score:     2.1567                Validation Mean MAE Score:
2.2642
Train Std MAE Score:      0.0529                Validation Std MAE Score:
0.35
```

```
Train Mean RMSE Score:    10.5865               Validation Mean RMSE Score:
11.1579
Train Std RMSE Score:     0.4678                Validation Std RMSE Score:
3.5502
```

## 3.6 Neural Network regression

### 3.6.1 Default is 1 hidden layer with 100 neurons

```
[ ]: from sklearn.neural_network import MLPRegressor
     train_scores, val_scores = cross_validation(MLPRegressor(random_state=1,
     ↪max_iter=500), X, y, 10)
```

```
*****
Model 1 trained
Training Scores
R2 Score : 0.8442          MAE : 2.6132      RMSE : 11.8528
Validation Scores
R2 Score : 0.844          MAE : 2.8365      RMSE : 10.7524
*****
Model 2 trained
Training Scores
R2 Score : 0.8553          MAE : 2.0893      RMSE : 11.2233
Validation Scores
R2 Score : 0.8328          MAE : 2.4967      RMSE : 13.0572
*****
```

```

Model 3 trained
Training Scores
R2 Score : 0.8337      MAE : 2.6943      RMSE : 11.9967
Validation Scores
R2 Score : 0.9152      MAE : 2.7922      RMSE : 9.4942
*****
Model 4 trained
Training Scores
R2 Score : 0.8565      MAE : 2.3367      RMSE : 11.2896
Validation Scores
R2 Score : 0.8268      MAE : 2.4525      RMSE : 12.2233

/usr/local/lib/python3.7/dist-
packages/sklearn/neural_network/_multilayer_perceptron.py:571:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (500) reached and
the optimization hasn't converged yet.
  % self.max_iter, ConvergenceWarning)

*****
Model 5 trained
Training Scores
R2 Score : 0.8811      MAE : 2.6773      RMSE : 10.3188
Validation Scores
R2 Score : 0.6288      MAE : 3.0712      RMSE : 17.177
*****
Model 6 trained
Training Scores
R2 Score : 0.8545      MAE : 2.7583      RMSE : 11.0693
Validation Scores
R2 Score : 0.7661      MAE : 3.2038      RMSE : 17.2891
*****
Model 7 trained
Training Scores
R2 Score : 0.8517      MAE : 2.6029      RMSE : 11.8045
Validation Scores
R2 Score : 0.8762      MAE : 2.3416      RMSE : 7.004
*****
Model 8 trained
Training Scores
R2 Score : 0.8421      MAE : 3.4655      RMSE : 11.8719
Validation Scores
R2 Score : 0.8178      MAE : 3.0049      RMSE : 12.2158
*****
Model 9 trained
Training Scores
R2 Score : 0.839       MAE : 2.7145      RMSE : 12.0548
Validation Scores
R2 Score : 0.9034      MAE : 2.6728      RMSE : 8.4049
*****

```

```

Model 10 trained
Training Scores
R2 Score : 0.8261      MAE : 2.698      RMSE : 12.2111
Validation Scores
R2 Score : 0.971      MAE : 2.6819     RMSE : 5.741

```

```
[ ]: print("Neural network regression results")
      print_results(train_scores, val_scores)
```

Neural network regression results

Mean and standard deviation of 10-fold cross validation results

```

-----
Train Mean R2 Score:      0.8484      Validation Mean R2 Score:
0.8382
Train Std R2 Score:      0.0145      Validation Std R2 Score:
0.0888

Train Mean MAE Score:     2.665      Validation Mean MAE Score:      2.7554
Train Std MAE Score:      0.3316      Validation Std MAE Score:
0.267

Train Mean RMSE Score:    11.5693     Validation Mean RMSE Score:
11.3359
Train Std RMSE Score:     0.554      Validation Std RMSE Score:      3.6943

```

### 3.7 Neural Network regression

#### 3.7.1 2 hidden layer with 64 neurons and 32 neurons

```
[70]: from sklearn.neural_network import MLPRegressor
      train_scores, val_scores = _
      ↪cross_validation(MLPRegressor(hidden_layer_sizes=(64, 64), random_state=1,_
      ↪max_iter=500), X, y, 10)
```

```

*****
Model 1 trained
Training Scores
R2 Score : 0.8537      MAE : 2.4604      RMSE : 11.4837
Validation Scores
R2 Score : 0.8638      MAE : 2.7862      RMSE : 10.0478
*****
Model 2 trained
Training Scores
R2 Score : 0.8425      MAE : 3.1197      RMSE : 11.7073

```

```

Validation Scores
R2 Score : 0.8267      MAE : 3.4537      RMSE : 13.2953
*****
Model 3 trained
Training Scores
R2 Score : 0.8329      MAE : 2.6483      RMSE : 12.0261
Validation Scores
R2 Score : 0.9025      MAE : 2.8511      RMSE : 10.1812
*****
Model 4 trained
Training Scores
R2 Score : 0.8446      MAE : 3.3345      RMSE : 11.7464
Validation Scores
R2 Score : 0.8172      MAE : 3.3833      RMSE : 12.5576
*****
Model 5 trained
Training Scores
R2 Score : 0.8849      MAE : 2.3674      RMSE : 10.1522
Validation Scores
R2 Score : 0.592       MAE : 2.8878      RMSE : 18.0079
*****
Model 6 trained
Training Scores
R2 Score : 0.8635      MAE : 2.456       RMSE : 10.7184
Validation Scores
R2 Score : 0.7768      MAE : 3.0608      RMSE : 16.889
*****
Model 7 trained
Training Scores
R2 Score : 0.8505      MAE : 2.8334      RMSE : 11.8505
Validation Scores
R2 Score : 0.8768      MAE : 2.6092      RMSE : 6.9869
*****
Model 8 trained
Training Scores
R2 Score : 0.8533      MAE : 2.7682      RMSE : 11.4416
Validation Scores
R2 Score : 0.8308      MAE : 2.4281      RMSE : 11.7737
*****
Model 9 trained
Training Scores
R2 Score : 0.8372      MAE : 3.972       RMSE : 12.1236
Validation Scores
R2 Score : 0.8973      MAE : 3.7011      RMSE : 8.6678
*****
Model 10 trained
Training Scores
R2 Score : 0.833       MAE : 3.017       RMSE : 11.9659

```



Validation Scores

R2 Score : 0.9709                      MAE : 2.8536              RMSE : 5.7455

```
[71]: print("Neural network regression results")
      print_results(train_scores, val_scores)
```

Neural network regression results

Mean and standard deviation of 10-fold cross validation results

```
-----
Train Mean R2 Score:      0.8496                Validation Mean R2 Score:
0.8355
Train Std R2 Score:      0.015                Validation Std R2 Score:      0.0961

Train Mean MAE Score:    2.8977                Validation Mean MAE Score:
3.0015
Train Std MAE Score:    0.4647                Validation Std MAE Score:
0.3781

Train Mean RMSE Score:   11.5216                Validation Mean RMSE Score:
11.4153
Train Std RMSE Score:    0.5946                Validation Std RMSE Score:
```